

Sparkify

May 27, 2021

1 Sparkify Project Workspace

This workspace contains a tiny subset (128MB) of the full dataset available (12GB). Feel free to use this workspace to build your project, or to explore a smaller subset with Spark before deploying your cluster on the cloud. Instructions for setting up your Spark cluster is included in the last lesson of the Extracurricular Spark Course content.

You can follow the steps below to guide your data analysis and model building portion of this project.

```
In [112]: # import libraries
          from pyspark.sql import SparkSession, SQLContext
          from pyspark.sql.functions import avg, col, concat, desc, explode, lit, min, max, split
          from pyspark.sql.types import IntegerType
          from pyspark.ml.classification import LogisticRegression, RandomForestClassifier, GBTC
          from pyspark.ml.feature import StandardScaler, RegexTokenizer, StringIndexer, CountVect
          from pyspark.ml import Pipeline
          from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
          from pyspark.ml.evaluation import BinaryClassificationEvaluator, MulticlassClassifica

          import seaborn as sns
          import matplotlib.pyplot as plt
          %matplotlib inline

          import numpy as np
          import pandas as pd

In [113]: # create a Spark session
          spark = SparkSession.builder.master("local").appName("Capstone_Project").getOrCreate()
```

2 Load and Clean Dataset

In this workspace, the mini-dataset file is `mini_sparkify_event_data.json`. Load and clean the dataset, checking for invalid or missing data - for example, records without userids or sessionids.

```
In [114]: # load data into spark DataFrame

          mydata = spark.read.json("./mini_sparkify_event_data.json")
```

```
mydata.printSchema()
```

```
root
|-- artist: string (nullable = true)
|-- auth: string (nullable = true)
|-- firstName: string (nullable = true)
|-- gender: string (nullable = true)
|-- itemInSession: long (nullable = true)
|-- lastName: string (nullable = true)
|-- length: double (nullable = true)
|-- level: string (nullable = true)
|-- location: string (nullable = true)
|-- method: string (nullable = true)
|-- page: string (nullable = true)
|-- registration: long (nullable = true)
|-- sessionId: long (nullable = true)
|-- song: string (nullable = true)
|-- status: long (nullable = true)
|-- ts: long (nullable = true)
|-- userAgent: string (nullable = true)
|-- userId: string (nullable = true)
```

2.1 Clean Data

```
In [115]: #Clean Dataset
```

```
# temp view of the data frame
```

```
mydata.createOrReplaceTempView('data_tbl')
```

```
In [116]: # check if there are nulls in sessionId column
```

```
spark.sql("""
    SELECT COUNT(userId) as UserId
    FROM data_tbl
    WHERE sessionId IS NULL
    """).show()
```

```
+-----+
|UserId|
+-----+
|      0|
+-----+
```

```
In [117]: # check if there are empty sessionIds
```

```
spark.sql("""
    SELECT COUNT(userId) as UserId
    FROM data_tbl
    WHERE sessionId == ''
    """).show()
```

```
+-----+
|UserId|
+-----+
|      0|
+-----+
```

```
In [118]: # check if there are nulls in userId column
```

```
spark.sql("""
    SELECT COUNT(userId) as UserId
    FROM data_tbl
    WHERE userId IS NULL
    """).show()
```

```
+-----+
|UserId|
+-----+
|      0|
+-----+
```

```
In [119]: # check if there are empty UserIDs
```

```
spark.sql("""
    SELECT COUNT(userId) as UserId
    FROM data_tbl
    WHERE userId == ''
    """).show()
```

```
+-----+
|UserId|
+-----+
|  8346|
+-----+
```

```
In [120]: # remove the invalid user IDs from the dataset
```

```
mydata = spark.sql("""
    SELECT *
    FROM data_tbl
    WHERE userId != ''
    """)
```

In [121]: *# temporary view of the data frame*

```
mydata.createOrReplaceTempView('data_tbl')
```

3 Exploratory Data Analysis

When you're working with the full dataset, perform EDA by loading a small subset of the data and doing basic manipulations within Spark. In this workspace, you are already provided a small subset of data you can explore.

3.0.1 Define Churn

Once you've done some preliminary analysis, create a column Churn to use as the label for your model. I suggest using the Cancellation Confirmation events to define your churn, which happen for both paid and free users. As a bonus task, you can also look into the Downgrade events.

3.0.2 Explore Data

Once you've defined churn, perform some exploratory data analysis to observe the behavior for users who stayed vs users who churned. You can start by exploring aggregates on these two groups of users, observing how much of a specific action they experienced per a certain time unit or number of songs played.

In [122]: `page = mydata.select("page").dropDuplicates().show()`

```
+-----+
|           page|
+-----+
|           Cancel|
| Submit Downgrade|
|           Thumbs Down|
|               Home|
|           Downgrade|
|           Roll Advert|
|               Logout|
|           Save Settings|
|Cancellation Conf...|
|               About|
|               Settings|
| Add to Playlist|
|           Add Friend|
|           NextSong|
```

```
|           Thumbs Up|
|           Help|
|           Upgrade|
|           Error|
|       Submit Upgrade|
+-----+
```

In [123]: *# create churn user list*

```
mydata = spark.sql("""
    SELECT *,
    CASE
        WHEN page == 'Cancellation Confirmation' THEN 1
        ELSE 0 END as Churned
    FROM data_tbl
""")

mydata.createOrReplaceTempView('data_tbl')

Churned = spark.sql("""
    SELECT DISTINCT userID
    FROM data_tbl
    WHERE Churned = 1
""").toPandas().values

Churned = [user[0] for user in Churned]
```

In [124]: *#show churned and non-churned user in dataset*

```
spark.sql("""
    SELECT
        Churned,
        count(distinct userID)
    FROM
        data_tbl
    GROUP BY
        Churned
""")
```

Out[124]: DataFrame[Churned: int, count(DISTINCT userID): bigint]

In [125]: *#create churn table*

```
churn = spark.sql("""
    SELECT
        distinct userID,
        Churned
```

```

        FROM
            data_tbl

        """)
    churn.createOrReplaceTempView('churn')

In [126]: # show churn in gender

    spark.sql("""
        SELECT distinct
            gender,
            Churned,
            count(distinct userId) as DistinctUsers
        FROM
            data_tbl
        GROUP BY
            gender, Churned
        order by Churned desc
        """)

Out[126]: DataFrame[gender: string, Churned: int, DistinctUsers: bigint]

```

4 Feature Engineering

Once you've familiarized yourself with the data, build out the features you find promising to train your model on. To work with the full dataset, you can follow the following steps. - Write a script to extract the necessary features from the smaller subset of data - Ensure that your script is scalable, using the best practices discussed in Lesson 3 - Try your script on the full data set, debugging your script if necessary

If you are working in the classroom workspace, you can just extract features based on the small subset of data contained here. Be sure to transfer over this work to the larger dataset when you work on your Spark cluster.

```

In [127]: #gender feature " replace str by int values"

    gender = mydata.dropDuplicates(['userId']).sort('userId').select(['userId', 'gender'])
    gender = gender.replace(['F', 'M'], ['1', '0'], 'gender')
    gender = gender.withColumn('gender', gender.gender.cast("int"))

    gender.createOrReplaceTempView('gender')

In [128]: #number of songs played per user

    songs = mydata.where(mydata.song!='null').groupby('userId')
    songs= songs.agg(count(mydata.song).alias('Played_Songs')).orderBy('userId')
    songs = songs.select(['userId', 'Played_Songs'])

    songs.createOrReplaceTempView('songs')

```

```
In [129]: # number of listened singers per user
```

```
listened_singers_per_user = mydata.dropDuplicates(['userId', 'artist']).groupBy('userId')
listened_singers_per_user = listened_singers_per_user.agg(count(mydata.artist).alias(''))
listened_singers_per_user = listened_singers_per_user.select(['userId', 'Listened_Singers'])

listened_singers_per_user.createOrReplaceTempView('listened_singers_per_user')
```

```
In [130]: #thumbs_Down
```

```
thumbs_Down = mydata.where(mydata.page=='Thumbs Down').groupBy(['userId'])
thumbs_Down = thumbs_Down.agg(count(col('page')).alias('thumbs_down')).orderBy('userId')
thumbs_Down = thumbs_Down.select(['userId', 'thumbs_down'])

thumbs_Down.createOrReplaceTempView('thumbs_Down')
```

```
In [131]: #thumbs_Up
```

```
thumbs_Up = mydata.where(mydata.page=='Thumbs Up').groupBy(['userId'])
thumbs_Up = thumbs_Up.agg(count(col('page')).alias('thumbs_Up')).orderBy('userId')
thumbs_Up = thumbs_Up.select(['userId', 'thumbs_Up'])

thumbs_Up.createOrReplaceTempView('thumbs_Up')
```

5 Modeling

Split the full dataset into train, test, and validation sets. Test out several of the machine learning methods you learned. Evaluate the accuracy of the various models, tuning parameters as necessary. Determine your winning model based on test accuracy and report results on the validation set. Since the churned users are a fairly small subset, I suggest using F1 score as the metric to optimize.

```
In [132]: # join features
```

```
Data = churn.dropDuplicates(['userId']).sort('userId').select(['userId', 'Churned'])
for selected_features in [gender, songs, listened_singers_per_user, thumbs_Up, thumbs_Down]:
    Data = Data.join(selected_features, 'userId')
```

```
In [133]: # convert data type into float
```

```
for selected_features in Data.columns[1:]:
    Data = Data.withColumn(selected_features, Data[selected_features].cast('float'))
```

```
In [134]: Data.dtypes
```

```
Out[134]: [('userId', 'string'),
           ('Churned', 'float'),
           ('gender', 'float'),
           ('Played_Songs', 'float'),
```

```

('Listened_Singers', 'float'),
('thumbs_Up', 'float'),
('thumbs_down', 'float')]

```

In [135]: *# split our data into train and test sets*

```

train_set, test_set = Data.randomSplit([0.8, 0.2])

```

In [136]: `assembler = VectorAssembler(inputCols=Data.columns[2:], outputCol='featuresassemble')`
`scaler = StandardScaler(inputCol="featuresassemble", outputCol="features")`
`indexer = StringIndexer(inputCol="Churned", outputCol="label")`
`stringIndexer = StringIndexer(inputCol="label", outputCol="indexed")`
`RandomForestClassifier = RandomForestClassifier(numTrees=3, maxDepth=2, labelCol="indexed")`
`LogisticRegression = LogisticRegression(maxIter=100, regParam=0.0, elasticNetParam=0)`

In [137]: `LogisticRegression_pipeline = Pipeline(stages=[assembler, scaler, indexer, LogisticRegression])`

```

paramGrid_LogisticRegression = ParamGridBuilder().addGrid(LogisticRegression.regParam,

```

```

CrossValidator_LogisticRegression = CrossValidator(estimator=LogisticRegression_pipeline,
                                                    evaluator=BinaryClassificationEvaluator(), numFolds=5)

```

```

CrossValidator_LogisticRegression_Model = CrossValidator_LogisticRegression.fit(train_set)

```

```

CrossValidator_LogisticRegression_Model.avgMetrics

```

Out[137]: [0.6658790392910175, 0.7057689613270375, 0.7096003789515586]

In [138]: `RandomForest_pipeline = Pipeline(stages=[assembler, scaler, indexer, stringIndexer, RandomForestClassifier])`

```

paramGrid_RandomForest = ParamGridBuilder().addGrid(RandomForestClassifier.numTrees, [1, 3, 10, 20, 50, 100])

```

```

CrossValidator_RandomForest = CrossValidator(estimator=RandomForest_pipeline, estimatorParamMaps=paramGrid_RandomForest,
                                              evaluator=BinaryClassificationEvaluator(), numFolds=3)

```

```

CrossValidator_RandomForest_Model = CrossValidator_RandomForest.fit(train_set)

```

```

CrossValidator_RandomForest_Model.avgMetrics

```

Out[138]: [0.6497828459602506, 0.6037712161804721]

6 Final Steps

Clean up your code, adding comments and renaming variables to make the code easier to read and maintain. Refer to the Spark Project Overview page and Data Scientist Capstone Project Rubric to make sure you are including all components of the capstone project and meet all expectations. Remember, this includes thorough documentation in a README file in a Github repository, as well as a web app or blog post.


```
In [ ]:
```